

predicting_profitability_movies

February 11, 2019

Alejandro Garcia
Predicting profitability from movies dataset

0.1 Table of Content

0.1.1 1.) Introduction

0.1.2 2.) Dataset

0.1.3 3.) Exploratory Data Analysis

0.1.4 4.) Data Wrangling

0.1.5 5.) Machine Learning

0.1.6 6.) Conclusion

###

Introduction

The main objective of this project is to accurately predict the profitability of a movie based on numerous parameters shown below. I will start with EDA to get familiar with the data and start to think about crafting a possible solution to this problem. I will then proceed to wrangling and cleaning the data for our model to work properly. Finally, I will configure several machine learning models and decide which one gives us the best results.

Note: If you have any questions about this project please feel free to shoot over an email: a19garcia95@gmail.com.

###

Dataset

The data comes from www.IMBd.com which is a website used to review movies. Here is a preview of what the data looks like.

```
In [3]: import pandas as pd
df = pd.read_csv("box_office_predictions.csv")
df.head()
```

```
Out [3]:
```

	budget	country	director	genre	gross	\
0	237000000.0	UK	James Cameron	Action	760507625.0	
1	200000000.0	USA	James Cameron	Drama	658672302.0	
2	150000000.0	USA	Colin Trevorrow	Action	652270625.0	
3	220000000.0	USA	Joss Whedon	Action	623357910.0	

```
4 185000000.0      USA  Christopher Nolan  Action  534858444.0
```

	name	rating	runtime	score	star	\
0	Avatar (2009)	PG-13	162	7.8	Sam Worthington	
1	Titanic (1997)	PG-13	194	7.8	Leonardo DiCaprio	
2	Jurassic World (2015)	PG-13	124	7.0	Chris Pratt	
3	The Avengers (2012)	PG-13	143	8.1	Robert Downey Jr.	
4	The Dark Knight (2008)	PG-13	152	9.0	Christian Bale	

	studio	votes
0	Twentieth Century Fox Film Corporation	958400
1	Twentieth Century Fox Film Corporation	865551
2	Universal Pictures	470625
3	Marvel Studios	1069292
4	Warner Bros.	1845853

```
###
```

Exploratory Data Analysis

```
In [4]: df.describe()
```

```
Out [4]:
```

	budget	gross	runtime	score	votes
count	6.000000e+03	6.000000e+03	6000.000000	6000.000000	6.000000e+03
mean	2.469918e+07	3.341635e+07	106.587000	6.386383	7.188537e+04
std	3.721710e+07	5.735205e+07	18.026885	0.994921	1.308033e+05
min	0.000000e+00	4.410000e+02	50.000000	1.500000	2.700000e+01
25%	0.000000e+00	1.527796e+06	95.000000	5.800000	7.791750e+03
50%	1.100000e+07	1.229897e+07	102.000000	6.500000	2.660150e+04
75%	3.262500e+07	4.007256e+07	115.000000	7.100000	7.677475e+04
max	3.000000e+08	7.605076e+08	366.000000	9.300000	1.868308e+06

Interesting enough, when we run the describe() function through the dataset, we are able to observe that the upper quartile is where most of the data is concentrated. An example of this is the cell below that shows that the top 6 countries account for most of the countries where movies are made.

```
In [5]: df['country'].value_counts()
```

```
Out [5]: USA          4281
         UK           615
         France       249
         Canada       126
         Germany      119
         Australia     71
         Japan         59
         Spain         50
         Italy         47
         Ireland       40
         Hong Kong     38
```

India	36
Denmark	31
China	23
South Korea	18
New Zealand	17
Sweden	17
Belgium	15
Mexico	13
Netherlands	12
Russia	10
Iran	10
Argentina	10
Czech Republic	8
Norway	8
Taiwan	7
Hungary	7
Switzerland	6
South Africa	6
Brazil	5
West Germany	4
Greece	4
Israel	4
Austria	4
Romania	3
Chile	3
Finland	3
Peru	2
Colombia	2
Portugal	2
Indonesia	2
Poland	2
Thailand	2
Iceland	1
Bahamas	1
Republic of Macedonia	1
Malta	1
Palestine	1
Federal Republic of Yugoslavia	1
Jamaica	1
Saudi Arabia	1
Soviet Union	1

Name: country, dtype: int64

```
In [6]: df[['score', 'votes', 'budget', 'gross', 'runtime']].corr()
```

```
Out[6]:
```

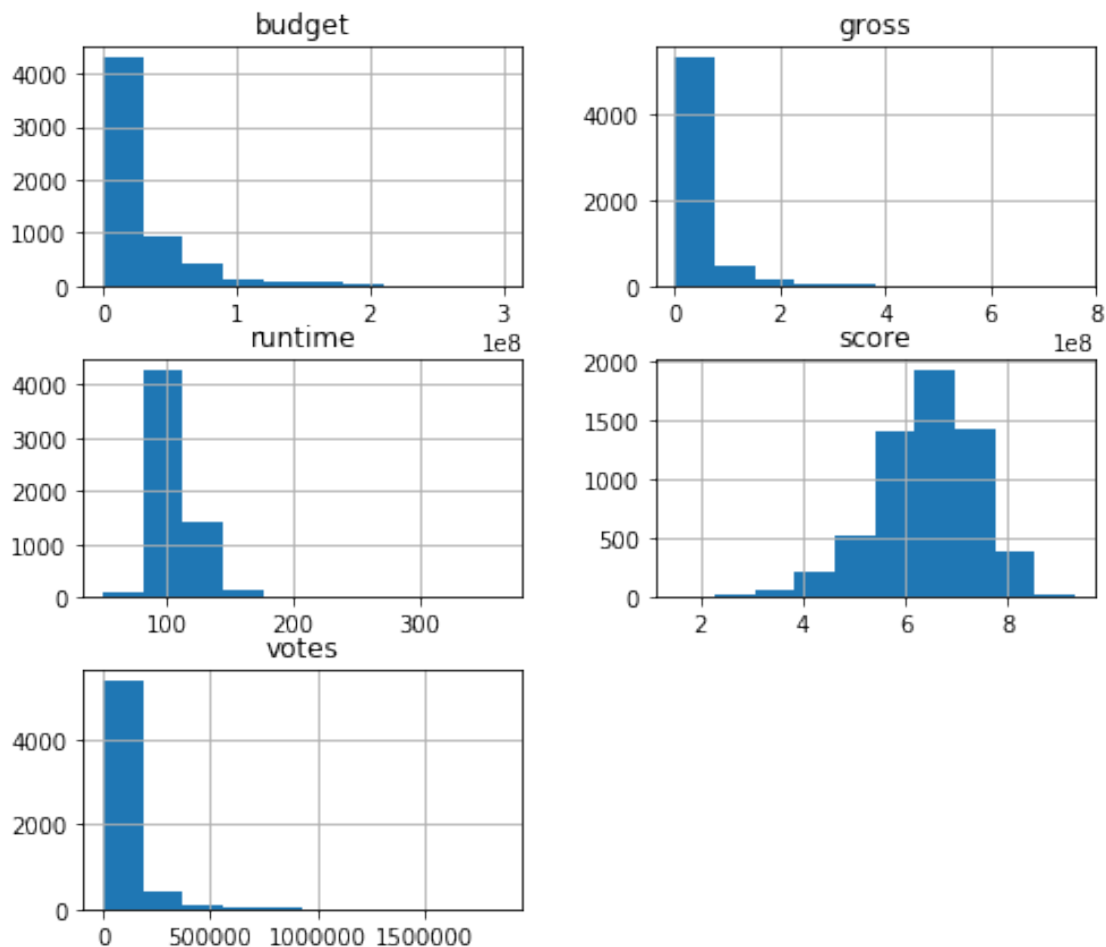
	score	votes	budget	gross	runtime
score	1.000000	0.393703	0.039169	0.161371	0.393470
votes	0.393703	1.000000	0.504808	0.664010	0.308524

budget	0.039169	0.504808	1.000000	0.716826	0.264963
gross	0.161371	0.664010	0.716826	1.000000	0.224996
runtime	0.393470	0.308524	0.264963	0.224996	1.000000

We can state that if the correlation between 2 variables is over 50%, it is highly correleated and therefore important for our machine learning model later on. In this case, budget-votes, gross-votes, budget-gross, are highly correlated. Later on I will run the same function on our cleaned dataset and observe wether the correleations remain the same.

Another graph that I find very helpful when trying to understand the data given to use, is the histogram. In this case we can observe how our numerical variables behave. (Histogram shown below)

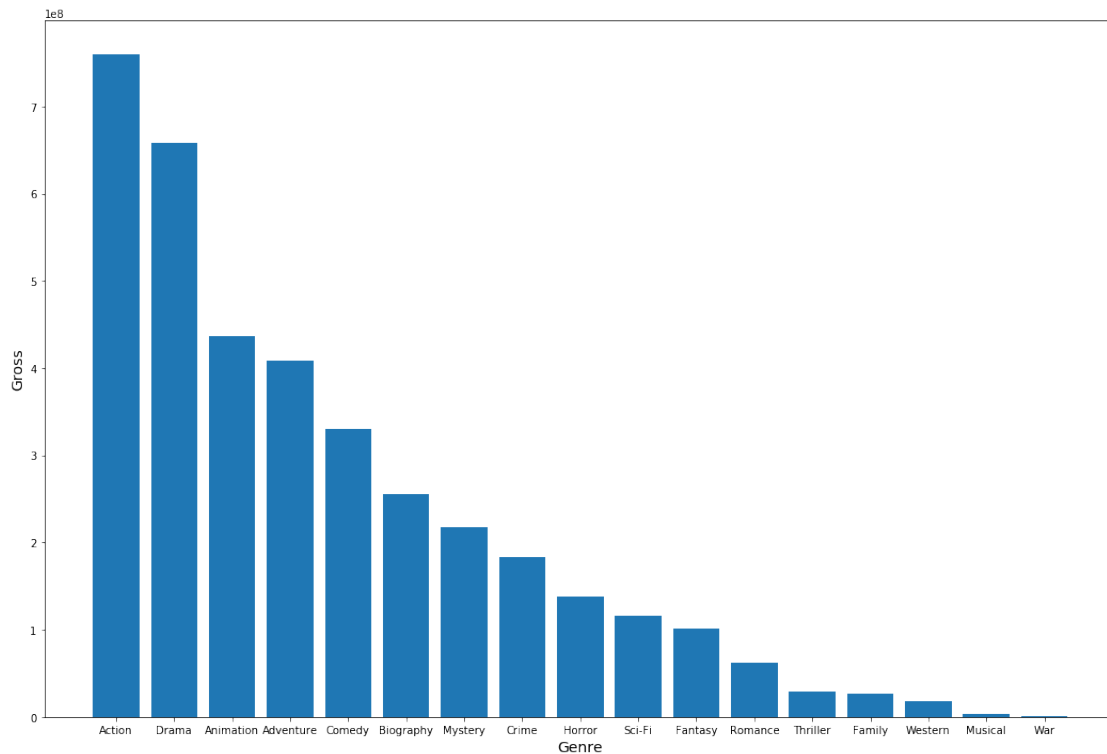
```
In [7]: import matplotlib.pyplot as plt
        %matplotlib inline
        df.hist(figsize=(8,7))
        plt.show()
```



0.1.7 NOTE: The next couple of bar graphs show how different categorical variables relate to our most important variable so far, which is 'Gross'. This is the revenue generated. I state that 'Gross' is our most important variable because the more revenue generated, the better the movie. This is generally how it goes. Of course, this is not always the case.

```
In [8]: x_axis = df['genre']
        y_axis = df['gross']
        width = 18
        height = 12
        plt.figure(figsize=(width, height))
        plt.bar(x_axis, y_axis)
        plt.xlabel('Genre', fontsize=14)
        plt.ylabel('Gross', fontsize=14)
```

Out[8]: Text(0,0.5,'Gross')



```
In [9]: print(df['genre'].head(10))
```

```
0    Action
1    Drama
2    Action
3    Action
4    Action
5    Action
```

```

6      Action
7      Action
8      Action
9      Animation
Name: genre, dtype: object

```

```

In [10]: xx = df['studio'].where(df['gross'] > 200000000)
          type(xx)

```

```

Out[10]: pandas.core.series.Series

```

```

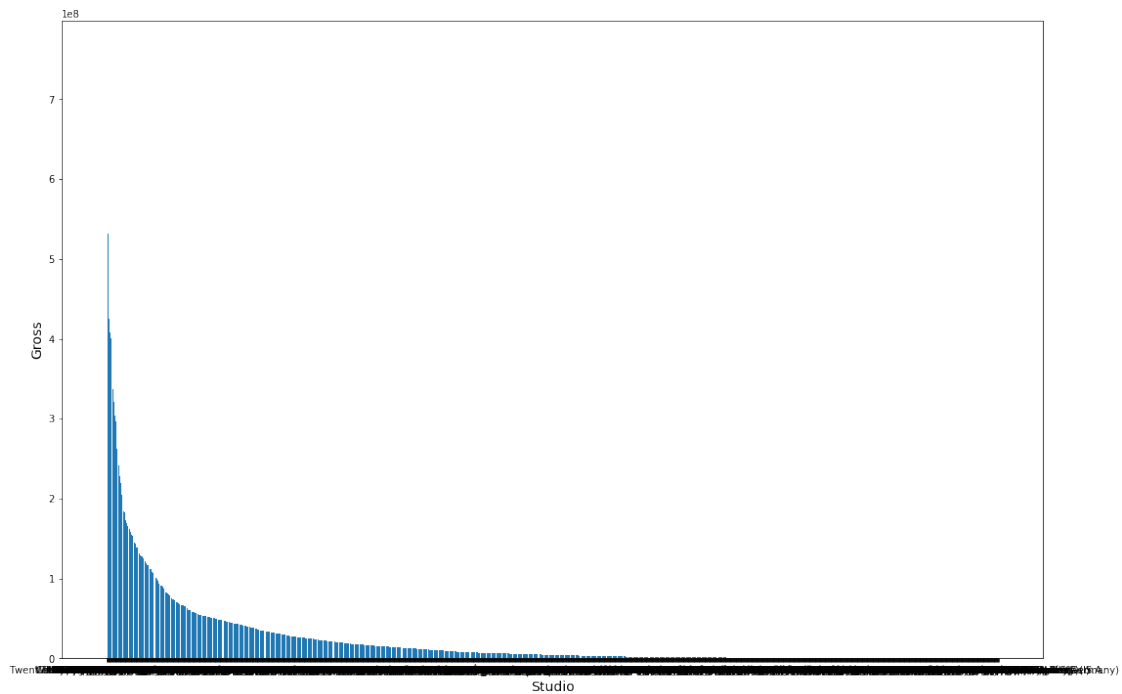
In [11]: x_axis = df['studio']
          y_axis = df['gross']
          width = 18
          height = 12
          plt.figure(figsize=(width, height))
          plt.bar(x_axis, y_axis)
          plt.xlabel('Studio', fontsize=14)
          plt.ylabel('Gross', fontsize=14)

```

```

Out[11]: Text(0,0.5,'Gross')

```



```

In [12]: x_axis = df['director']
          y_axis = df['gross']

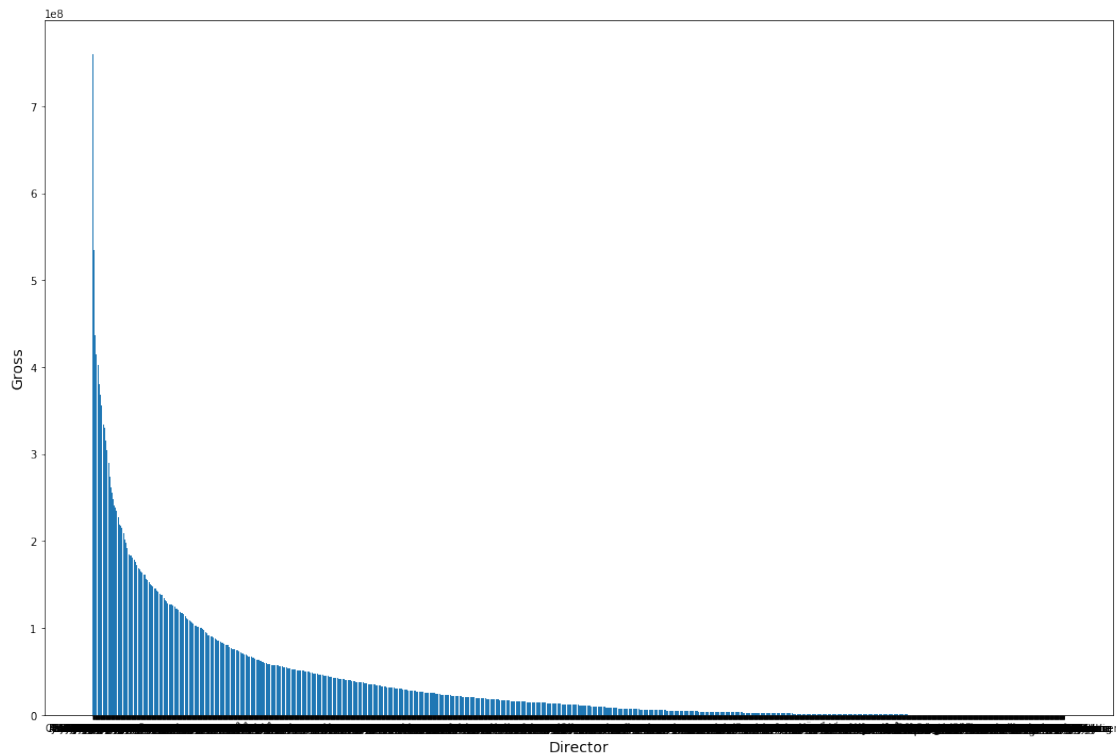
```

```

width = 18
height = 12
plt.figure(figsize=(width, height))
plt.bar(x_axis, y_axis)
plt.xlabel('Director', fontsize=14)
plt.ylabel('Gross', fontsize=14)

```

Out[12]: Text(0,0.5,'Gross')

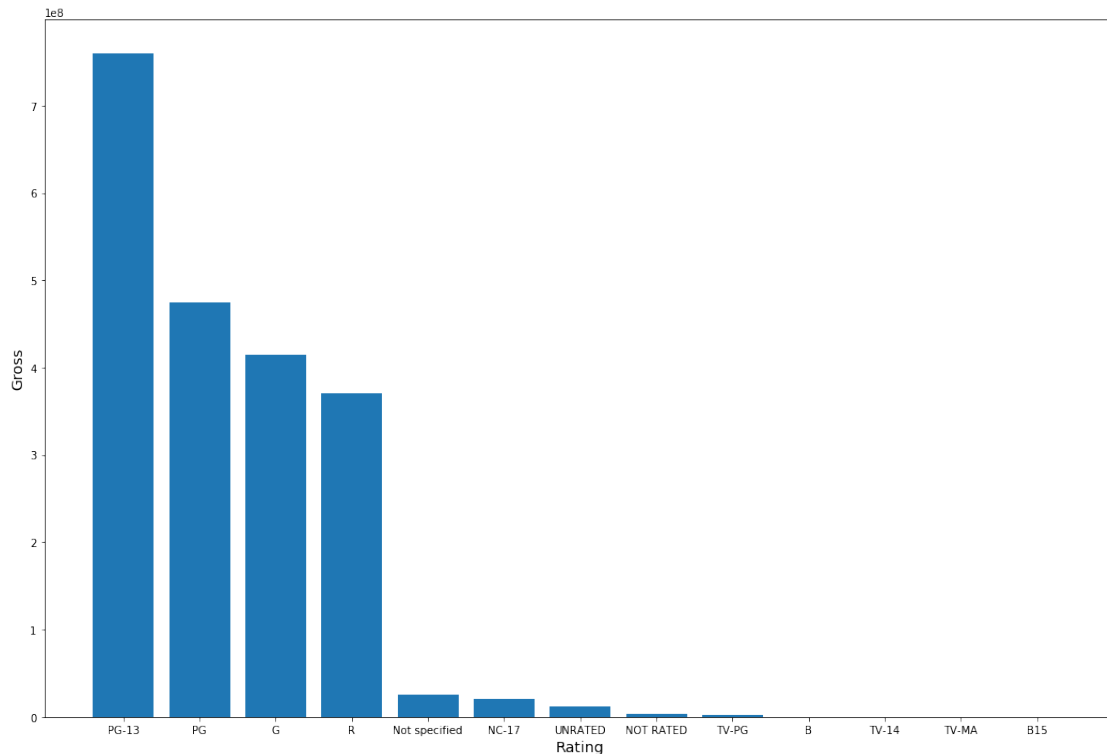


```

In [13]: x_axis = df['rating']
         y_axis = df['gross']
         width = 18
         height = 12
         plt.figure(figsize=(width, height))
         plt.bar(x_axis, y_axis)
         plt.xlabel('Rating', fontsize=14)
         plt.ylabel('Gross', fontsize=14)

```

Out[13]: Text(0,0.5,'Gross')



###

Data Wrangling

```
In [15]: df[df.budget == 0].head()
```

```
Out[15]:
```

	budget	country	director	genre	gross \
56	0.0	UK	David Yates	Adventure	295983305.0
207	0.0	USA	Walt Becker	Action	168273550.0
431	0.0	USA	John G. Avildsen	Action	115103979.0
553	0.0	USA	Nora Ephron	Comedy	95318203.0
592	0.0	USA	Tyler Perry	Comedy	90485233.0

	name	rating	runtime	score \
56	Harry Potter and the Deathly Hallows: Part 1 (...)	PG-13	146	7.7
207	Wild Hogs (2007)	PG-13	100	5.9
431	The Karate Kid Part II (1986)	PG	113	5.9
553	Michael (1996)	PG	105	5.7
592	Madea Goes to Jail (2009)	PG-13	103	4.3

	star	studio	votes
56	Daniel Radcliffe	Warner Bros.	370003
207	Tim Allen	Touchstone Pictures	104657
431	Pat Morita	Columbia Pictures Corporation	58596
553	John Travolta	Turner Pictures (I)	36553
592	Tyler Perry	Tyler Perry Company, The	10095


```
In [ ]: df = df.loc[df.budget > 0,:] #removing missing values
```

```
In [16]: #Profit would be a useful variable since not profitability is what determines success
df['profit'] = df.gross - df.budget
```

```
In [17]: df.describe(include=['object'])
```

```
Out[17]:
```

	country	director	genre	name	rating	\
count	6000	6000	6000	6000	6000	
unique	52	2549	17	6000	13	
top	USA	Woody Allen	Comedy	Gnomeo & Juliet (2011)	R	
freq	4281	29	1818	1	3009	

	star	studio
count	6000	6000
unique	2317	1996
top	Nicolas Cage	Universal Pictures
freq	38	269

```
In [18]: #lets condense the dataset, this will give us a sense of what matters the most
```

```
In [19]: studio_counts = df.studio.value_counts()
```

```
one_timers = studio_counts[studio_counts <= 1].index
three_timers = studio_counts[(studio_counts > 1) & (studio_counts <= 3)].index
five_timers = studio_counts[(studio_counts > 3) & (studio_counts <= 5)].index
ten_timers = studio_counts[(studio_counts > 5) & (studio_counts <= 10)].index
```

```
df['studio'].replace(one_timers, 'One Timer', inplace=True)
df['studio'].replace(three_timers, 'Three Timer', inplace=True)
df['studio'].replace(five_timers, 'Five Timer', inplace=True)
df['studio'].replace(ten_timers, 'Ten Timer', inplace=True)
```

```
In [20]: star_counts = df.star.value_counts()
```

```
one_timers = star_counts[star_counts <= 1].index
three_timers = star_counts[(star_counts > 1) & (star_counts <= 3)].index
five_timers = star_counts[(star_counts > 3) & (star_counts <= 5)].index
ten_timers = star_counts[(star_counts > 5) & (star_counts <= 10)].index
```

```
df['star'].replace(one_timers, 'One Timer', inplace=True)
df['star'].replace(three_timers, 'Three Timer', inplace=True)
df['star'].replace(five_timers, 'Five Timer', inplace=True)
df['star'].replace(ten_timers, 'Ten Timer', inplace=True)
```

```
In [21]: director_counts = df.director.value_counts()
```

```
one_timers = director_counts[director_counts <= 1].index
three_timers = director_counts[(director_counts > 1) & (director_counts <= 3)].index
```

```

five_timers = director_counts[(director_counts > 3) & (director_counts <= 5)].index
ten_timers = director_counts[(director_counts > 5) & (director_counts <= 10)].index

df['director'].replace(one_timers, 'One Timer', inplace=True)
df['director'].replace(three_timers, 'Three Timer', inplace=True)
df['director'].replace(five_timers, 'Five Timer', inplace=True)
df['director'].replace(ten_timers, 'Ten Timer', inplace=True)

```

```
In [22]: country_counts = df.country.value_counts()
```

```

other_countries = country_counts[country_counts < 50].index
df['country'].replace(other_countries, 'Other', inplace=True)

df.country.value_counts()

```

```

Out[22]: USA          4281
        UK           615
        Other        430
        France       249
        Canada       126
        Germany      119
        Australia     71
        Japan         59
        Spain         50
        Name: country, dtype: int64

```

```
In [23]: genre_counts = df.genre.value_counts()
```

```

other_genres = genre_counts[genre_counts < 50].index
df['genre'].replace(other_genres, 'Other', inplace=True)

df.genre.value_counts()

```

```

Out[23]: Comedy      1818
        Drama        1280
        Action       1175
        Crime         463
        Adventure     340
        Biography     309
        Animation     246
        Horror        243
        Other         126
        Name: genre, dtype: int64

```

```
In [24]: df['rating'].replace(['NOT RATED', 'UNRATED', 'Not specified'], 'NR', inplace=True)
```

```
In [57]: #Function will return how long the movie has been out
```

```
def extract_age(s, today=2014):
```

```

        return today - int( s[-5:-1] )

    extract_age('Titanic (1997)')

Out[57]: 17

In [58]: df['age'] = df.name.apply(extract_age)

###
Machine Learning

In [59]: #A good place to start diving into the ML section is by computing correlations

df[['score', 'votes', 'profit', 'budget']].corr()

Out[59]:
          score      votes      profit      budget
score    1.000000  0.393703  0.194065  0.039169
votes    0.393703  1.000000  0.480234  0.504808
profit   0.194065  0.480234  1.000000  0.096927
budget   0.039169  0.504808  0.096927  1.000000

In [60]: #We define as 'correlated' if its greater than 50%

In [61]: from sklearn.ensemble import RandomForestRegressor
         from sklearn.metrics import r2_score

In [62]: abt = pd.get_dummies ( df.drop(['name', 'gross', 'votes', 'score'], axis=1) ) #recall

In [63]: train = abt[abt.age >= 0]
         test = abt[abt.age <= 0]

         y_train = train.profit
         X_train = train.drop(['profit'], axis=1)

         y_test = test.profit
         X_test = test.drop(['profit'], axis=1)

In [64]: # Useful article for why choosing a random forest regressor: https://towardsdatascien

```

0.2 From the article above:

0.2.1 Random Forest is also considered as a very handy and easy to use algorithm, because it's default hyperparameters often produce a good prediction result. The number of hyperparameters is also not that high and they are straightforward to understand. One of the big problems in machine learning is overfitting, but most of the time this won't happen that easy to a random forest classifier. That's because if there are enough trees in the forest, the classifier won't overfit the model.

```

In [65]: # Train a basic random forest model
         rf = RandomForestRegressor(random_state=1234)
         rf.fit(X_train, y_train)

```

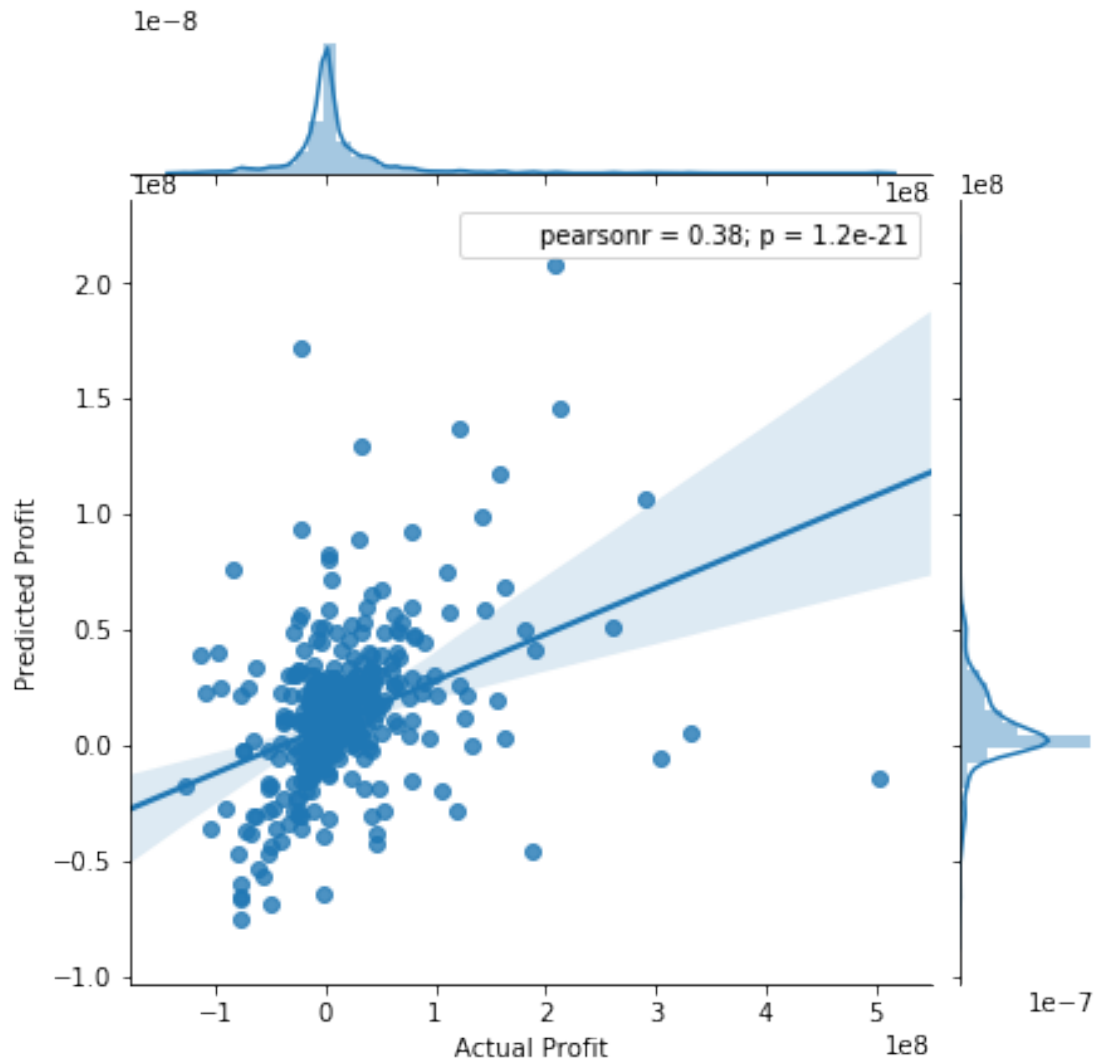
```
Out [65]: RandomForestRegressor(bootstrap=True, criterion='mse', max_depth=None,
                                max_features='auto', max_leaf_nodes=None,
                                min_impurity_decrease=0.0, min_impurity_split=None,
                                min_samples_leaf=1, min_samples_split=2,
                                min_weight_fraction_leaf=0.0, n_estimators=10, n_jobs=1,
                                oob_score=False, random_state=1234, verbose=0, warm_start=False)
```

```
In [66]: # Make prediction on test set
         pred = rf.predict(X_test)
```

```
In [68]: import seaborn as sns
```

```
In [69]: sns.jointplot(y_test, pred, kind='reg')
         plt.xlabel('Actual Profit')
         plt.ylabel('Predicted Profit')
         plt.show()
```

```
/Users/alejandrogarcia/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been
replaced by 'normalized'. Please see Matplotlib 1.2.0 for more information.
warnings.warn("The 'normed' kwarg is deprecated, and has been ")
/Users/alejandrogarcia/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been
replaced by 'normalized'. Please see Matplotlib 1.2.0 for more information.
warnings.warn("The 'normed' kwarg is deprecated, and has been ")
```



R² (coefficient of determination) regression score function.

Best possible score is 1.0 and it can be negative (because the model can be arbitrarily worse). A constant model that always predicts the expected value of y , disregarding the input features, would get a R² score of 0.0.

```
In [70]: r2_score(y_test, pred)
```

```
Out[70]: 0.12471120486001908
```

```
In [71]: # NOT A GREAT R2 SCORE
```

```
In [72]: abt_ps = pd.get_dummies ( df.drop(['name', 'gross', 'votes'], axis=1) ) #we kept the
```

```

train = abt_ps[abt_ps.age >= 0]
test = abt_ps[abt_ps.age <= 0]

y_train = train.profit
X_train = train.drop(['profit'], axis=1)

y_test = test.profit
X_test = test.drop(['profit'], axis=1)

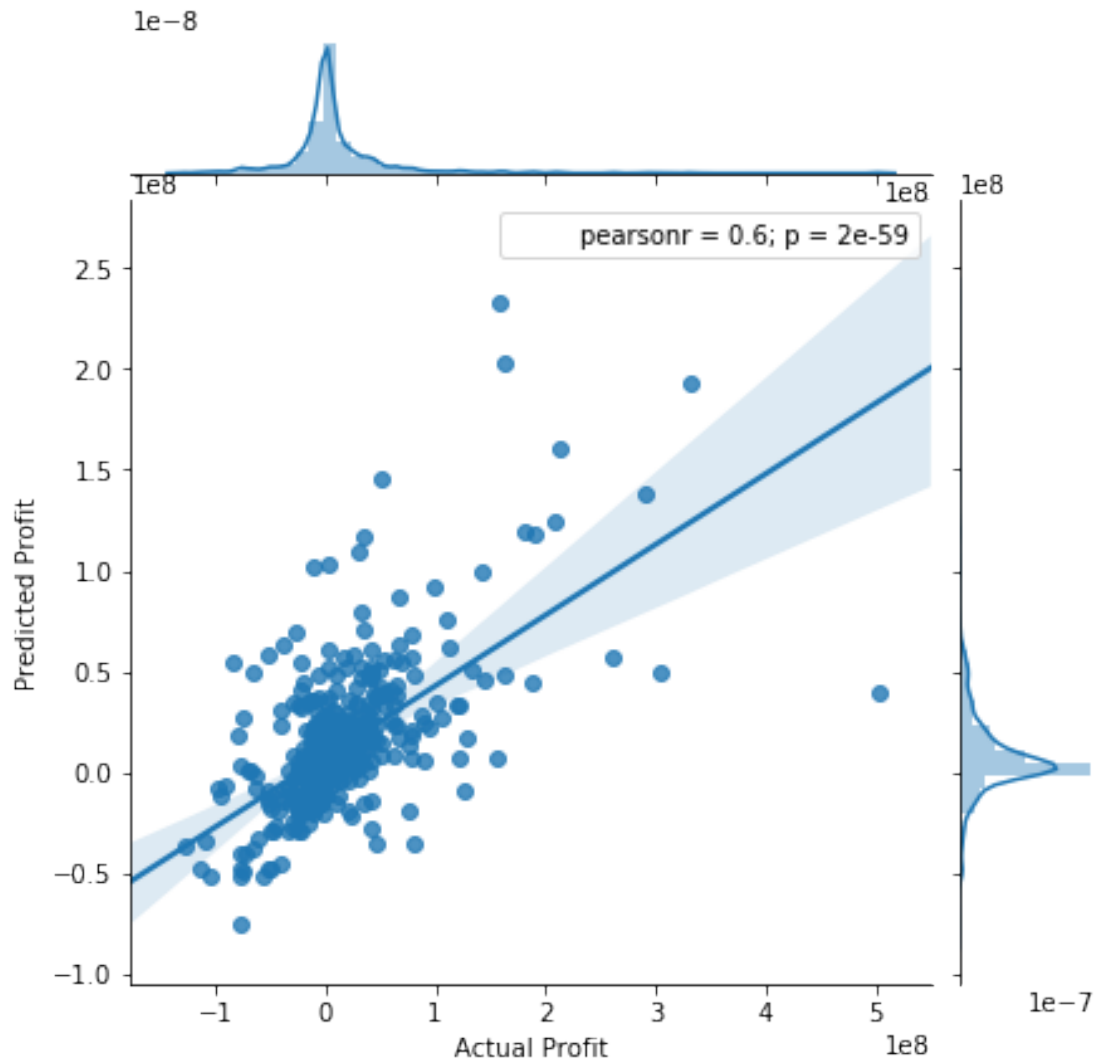
rf = RandomForestRegressor(random_state=1234)
rf.fit(X_train, y_train)

pred = rf.predict(X_test)

In [73]: sns.jointplot(y_test, pred, kind='reg')
plt.xlabel('Actual Profit')
plt.ylabel('Predicted Profit')
plt.show()

/Users/alejandrogarcia/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been
/Users/alejandrogarcia/anaconda3/lib/python3.6/site-packages/matplotlib/axes/_axes.py:6462: UserWarning: The 'normed' kwarg is deprecated, and has been

```



```
In [74]: r2_score(y_test, pred)
```

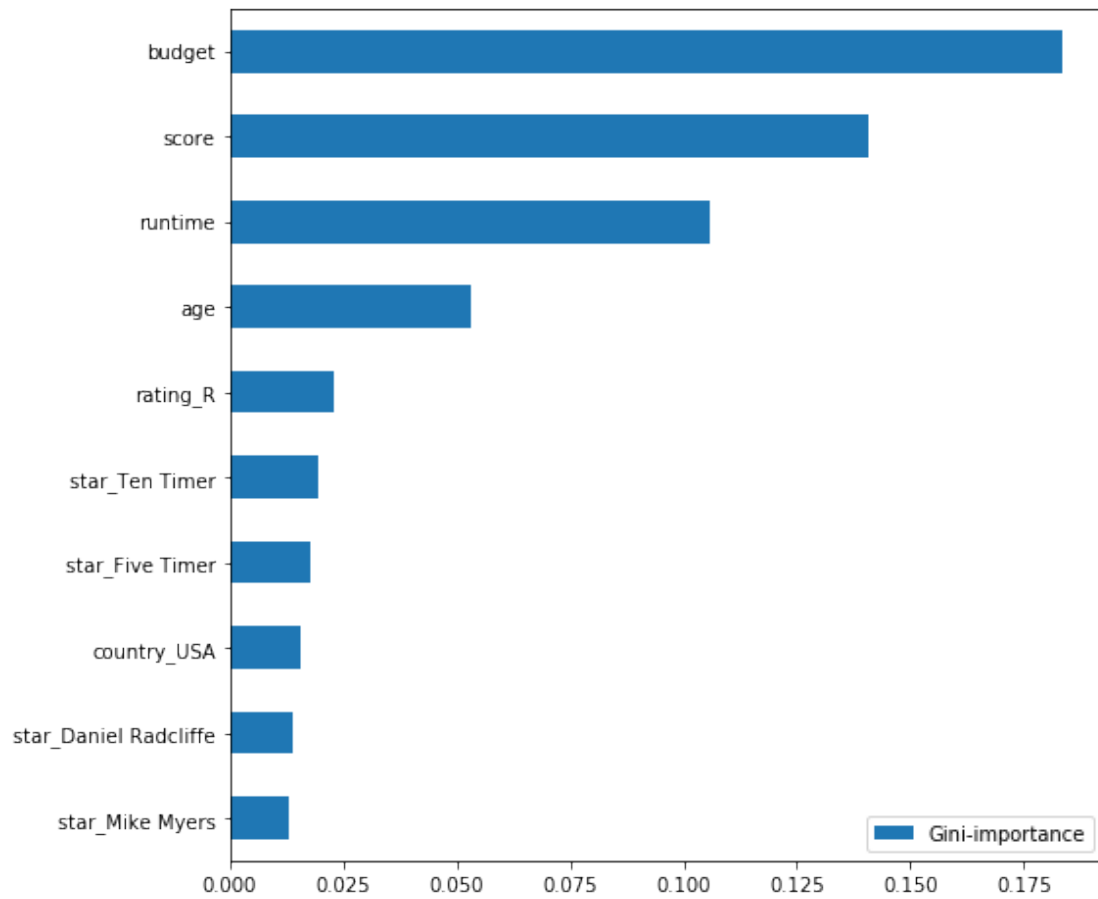
```
Out[74]: 0.36381519145681196
```

```
In [75]: #Much better R2 score
```

```
In [78]: def plot_feature_importances(columns, feature_importances, show_top_n=10):
    feats = dict( zip(columns, feature_importances) )
    imp = pd.DataFrame.from_dict(feats, orient='index').rename(columns={0: 'Gini-importance'})
    imp.sort_values(by='Gini-importance').tail(show_top_n).plot(kind='barh', figsize=(10, 10))
    plt.show()
```

```
In [79]: import matplotlib.pyplot as plt
    %matplotlib inline
```

```
plot_feature_importances(X_train.columns, rf.feature_importances_)
```



In [80]: *#As some Data Scientists might say, more data does not mean more results. Clearly only,*